



FAKULTÄT FÜR INFORMATIK

DER TECHNISCHEN UNIVERSITÄT MÜNCHEN

Masterarbeit in Informatik

to do title eng

Matthias Konstantin Fischer





FAKULTÄT FÜR INFORMATIK

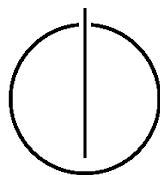
DER TECHNISCHEN UNIVERSITÄT MÜNCHEN

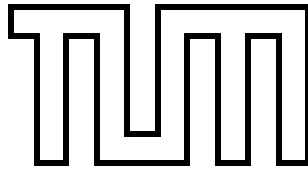
Masterarbeit in Informatik

to do title eng

to do, title ger

Author: Matthias Konstantin Fischer
Supervisor: Prof. Dr. Claudia Eckert
Advisor: M.Sc. Paul Muntean
Date: X November, 2016





FAKULTÄT FÜR INFORMATIK

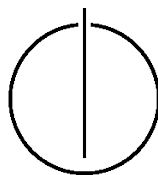
DER TECHNISCHEN UNIVERSITÄT MÜNCHEN

Masterarbeit in Informatik

to do title eng

to do, title ger

Author: Matthias Konstantin Fischer
Supervisor: Prof. Dr. Claudia Eckert
Advisor: M.Sc. Paul Muntean
Date: X November, 2016



Ich versichere, dass ich diese Diplomarbeit selbständig verfasst und nur die angegebenen Quellen und Hilfsmittel verwendet habe.

München, den 21. Oktober 2016

Matthias Konstantin Fischer

Acknowledgments

If someone contributed to the thesis... might be good to thank them here.

Abstract

An abstracts abstracts the thesis!

This is just an example that shows how long the abstract should be and which parts an abstract should contain.

Many applications such as the Chrome and Firefox browsers are largely implemented in C++ for its performance and modularity. Type casting, which converts one type of an object to another, plays an essential role in enabling polymorphism in C++ because it allows a program to utilize certain general or specific implementations in the class hierarchies. However, if not correctly used, it may return unsafe and incorrectly casted values, leading to so-called bad-casting or type-confusion vulnerabilities. Since a bad-casted pointer violates a programmer’s intended pointer semantics and enables an attacker to corrupt memory, bad-casting has critical security implications similar to those of other memory corruption vulnerabilities. Despite the increasing number of bad-casting vulnerabilities, the bad-casting detection problem has not been addressed by the security community.

In this paper, we present C A V E R , a runtime bad-casting detection tool. It performs program instrumentation at compile time and uses a new runtime type tracing mechanism—the type hierarchy table—to overcome the limitation of existing approaches and efficiently verify type casting dynamically. In particular, C A V E R can be easily and automatically adopted to target applications, achieves broader detection coverage, and incurs reasonable runtime overhead. We have applied C A V E R to large-scale software including Chrome and Firefox browsers, and discovered 11 previously unknown security vulnerabilities: nine in GNU libstdc++ and two in Firefox, all of which have been confirmed and subsequently fixed by vendors. Our evaluation showed that C A V E R imposes up to 7.6% and 64.6% overhead for performance-intensive benchmarks on the Chromium and Firefox browsers, respectively.

Contents

Acknowledgements	ix
Abstract	xi
1. Introduction	1
1.1. Motivation	2
1.2. Research Goals	2
1.3. Outline	2
2. Technical Overview	3
2.1. Types Inference/Recovery from Binaries	3
2.1.1. Types Recovery in General	3
2.1.2. Function Parameter Types Recovery	3
2.1.3. Caller/Callee Parameter Types Recovery	3
2.2. Code-Reuse Attacks	3
2.3. Control-Flow Integrity	3
2.4. Cntrol-Flow Integrity	3
I. Snippets	5
3. Snippet [AddressTaken]	7
3.1. step 1	7
3.2. step 2	7
4. Implementation	9
5. Evaluation	11
5.1. R1: Effectiveness of our Tool	11
5.2. R2: Precision of our Tool	15
5.3. R3: Performance overhead of our Tool	15
5.4. R4: Instrumentation overhead of our Tool	15
5.5. R5: Security coverage of our tool	15
5.6. R6: Which kind of attacks can our tool defend off	15
5.7. R7: Whar are the limitations of our Tool	15
5.8. R8: To Do.	15
5.9. Classification	15
5.9.1. Callsites	15
5.9.2. Calltargets	15

5.10. Patching Policies	15
5.10.1. AT	15
5.10.2. ParamCount	15
5.10.3. ParamType	15
5.11. Security Evaluation	16
5.12. Performance	16
6. Related Work	17
6.1. Mitigation of Advanced Code-Reuse Attacks	17
6.2. Type-Inference on Executables	18
6.3. Binary-based Protection against COOP	18
6.4. Source Code based Protection against COOP	18
6.5. Runtime-based Protection against COOP	19
7. Discussion	21
7.1. How to make the type inference more precise?	21
7.2. Comparison with TypeArmor and why are we better than TypeArmor? . . .	21
7.3. Whys is not TypeArmor working as it should to?	21
7.4. What is not clear bout TypeArmor?	21
7.5. What can for sure not work as in TypeArmor paper explained?	21
8. Conclusion and Future Work	23
8.1. Conclusion	23
8.2. Future Work	23
Bibliography	25

1. Introduction

Control-Flow Integrity (CFI) [32, 29] is one of the most used techniques to secure program execution flows against advanced Code-Reuse Attacks (CRAs).

Advanced CRAs such as COOP [3]

Proposal: Name for our tool:

CCTypeMapper (Caller Calle Type Mapper)

TypeShild

i have selected this name for our tool.

TypeProtector

CCTS (caller/calle type securer or shilder)

TypeFlower

you can also make your sugestion here.

Proposal:

Citation example: [?].

The introduction should answer this questions:

1.What is the problem?

Specify The Problem statement.

2-3 sentences.

2.What are the current solutions?

talk about TypeArmor [9]....

3.Where the solutions lack?

4.What is your idea?

5.Contributions.

In summary, we make the following contributions:

1. We did this

2. We did this

3. We did this.

The rest of the MA is organized as follows.

1.1. Motivation

here comes the Motivation.

2-3 sentences.

1.2. Research Goals

here comes the Motivation.

2-3 sentences.

1.3. Outline

here comes the Motivation.

2-3 sentences.

example:

The remainder of this thesis is organized as follows. Chapter 2 provides a profound background regarding VMs, VMI, and modern rootkits. We relate our work to previous research in Chapter 3. The design and architecture of WhiteRabbit is discussed in Chapter 4. This chapter comprises assumptions and necessary means that are required to meet the goals previously stated in Section 1.2. The WhiteRabbit prototype implementation is discussed and evaluated in Chapter 5 and Chapter 6, respectively. Finally, we provide an outlook concerning future work in Chapter 7 and conclude this thesis with a brief recapitulation in Chapter 8.

2. Technical Overview

2.1. Types Inference/Recovery from Binaries

2.1.1. Types Recovery in General

2.1.2. Function Parameter Types Recovery

2.1.3. Caller/Callee Parameter Types Recovery

2.2. Code-Reuse Attacks

this section will up to 1 DIN A page. First talk about code reuse attacks in general and then about COOP.

Say why COOP is not affected by the following mitigation approaches Talk about intel CET, talk about Windows CF guard.

2.3. Control-Flow Integrity

2.4. Cntrol-Flow Integrity

Part I.

Snippets

3. Snippet [AddressTaken]

As of now we used the full set of possible calltargets, which is the set of addresses of all function entry basic blocks. To further restrict the possible calltargets per callsite, we explored the notion of incorporating an address taken analysis into our application. The notion is that any indirect control flow instruction might only target addresses that are considered taken. An Address is considered to be a taken address, if it is loaded to memory or a register usually this is a constant, !optional! however it is also possible that simple calculations using multiplication and/or addition are used. We are not concerned with more complex calculations, because we have not observed compilers resorting to more complex methods and literature so far does agree [?].

Based on the notions of [?], introduced several types of indirect control flow targets of which only !shorthand! Code Pointer Constants (CK) and !shorthand! Computed Code Pointers (CC) are of interest to us. The reason for that is that the others are usually the target of indirect jumps, however we are (as of now) only interested in callsites.

!optional!

Definition 3.1 *!shorthand! Computed Code Pointers (CC) are, addresses that are computed during binary execution. In [?] this set only contains targets to intraprocedural indirect jumps and is thus of no interest for us*

Our approach of indentifying taken addresses is a two pronged approach. First, we iterate over the raw binary content of data segments additionally identifying possible dereferencable addresses. Second, we iterate over all instructions in functions within the disassembled binary.

3.1. step 1

We rely on Dyninst to tell us the boundaries of the .plt section !todo! add more information here and the .text section, which contain the executable part of the binary and thus are use to precheck any addresses that we might find in this step.

As suggested in [?], we slide a !todo!, how much byte ? 4 or 8 ? what happens on X86-64 compared to x86window over the data sections of the binary (namely the .data, the .rodata and the .dynsym).

3.2. step 2

We rely on Dyninst [?] to supply us with the correct function bounadries and addresses of instructions, which we then pass onto our instruction decoder, which is based on DynamoRIO. In essence there are types of analysis that are performed on each instruction. First we identify all relevant constants from the instruction

3. Snippet [AddressTaken]

1. If the instructions is a control flow instructions, we completely ignore it, as it cannot give us any information that is relevant. !todo! can we trace back from memory addresses and registers, what essentially is being called ?
2. We look a the sources and !todo! targets of the instructions and add it to the list of potentially interesting targets
3. If the target is a RIP-based address, we rely on DynamoRIO to decode it and also add it to the list of potentially interesting target
4. !todo! what is with constant functions ?
5. !todo! can we have DynamoRIO infer the result of simple lea instructions ?

Then for the resulting set of addresses, we check whether each either point to the entry block of a function, points within the .plt section, or is present in our reference map, which we calculated earlier. !todo! is the reference map needed ?

4. Implementation

In the Master thesis this section should be no longer than 1 DIN A page:

example of implementation text from USENIX cover Paper.

Please write in the same style.

We implemented C A V ER based on the LLVM Compiler project [43] (revision 212782, version 3.5.0). The static instrumentation module is implemented in Clang's CodeGen module and LLVM's Instrumentation module. The runtime library is implemented using the compiler-rt module based on LLVM's Sanitizer code base. In total, C A V ER is implemented in 3,540 lines of C++ code (excluding empty lines and comments). C A V ER is currently implemented for the Linux x86 platform, and there are a few platform-dependent mechanisms. For example, the type and tracing functions for global objects are placed in the .ctors section of ELF. As these platform-dependent features can also be found in other platforms, we believe C A V ER can be ported to other platforms as well. C A V ER interposes threading functions to maintain thread contexts and hold a per-thread red-black tree for stack objects. C A V ER also maintains the top and bottom addresses of stack segments to efficiently check pointer membership on the stack. We also modified the front-end drivers of Clang so that users of C A V ER can easily build and secure their target applications with one extra compilation flag and linker flag, respectively.

5. Evaluation

We evaluated our tool X with Y popular servers, by instrumenting them with our tool. We performed runtime performance test with the following applications.

Our Evaluation aims to answer the following research questions:

- **R1:** How effective is our tool in securing binary programs against the COOP attack?
- **R2:** How precise is our tool in detecting the types of the caller/caller pairs?
- **R3:** What is the performance overhead of our tool?
- **R4:** What are the instrumentation overheads imposed by our tool?
- **R5:** How many caller/called pairs are secured by our tool and how many remain unsecured?
- **R6:** Against which kind of attacks can our tool secure programs?
- **R7:** What are the Limitations of our Tool?
- **R8:** List is not exhaustive. Give another relevant research question. if there is one.

Comparison methods. Example: We used UBSAN (compare with TypeArmor), the state-of-art tool for detecting bad-casting bugs, as our comparison target of C A V E R . Also, We used C A V E R - NAIVE , which disabled the two optimization techniques described in §4.4, to show their effectiveness on runtime performance optimization.

Experimental setup. Example: All experiments were run on Ubuntu 13.10 (Linux Kernel 3.11) with a quad-core 3.40 GHz CPU (Intel Xeon E3-1245), 16 GB RAM, and 1 TB SSD-based storage.

5.1. R1: Effectiveness of our Tool

Table 5.1.: Classification CS

target	opt	#CS	problems	0	-1	-2	-3	-4	-5	-6	non-void-ok	non-void-probl.
x	x	x	x	x	x	x	x	x	x	x	x	x

Table 5.2.: Compound

opt	#CS	cs args (perfect %)	cs args (problem %)	cs non-void (correct %)	cs non-void (probl. %)	#ct	ct args (perfect %)	ct args (probl. %)	ct void (correct %)	ct void (correct %)
x	x	x	x	x	x	x	x	x	x	x

5. Evaluation

Table 5.3.: Classification CT

target	opt	#CS	problems	0	-1	-2	-3	-4	-5	-6	non-void-ok	non-void-probl.
x	x	x	x	x	x	x	x	x	x	x	x	x

Table 5.4.: Callsite Classification for paramcount

target	opt	problematic	+0	+1	+2	+3	+4	+5	+6
x	x	x	x	x	x	x	x	x	x

Table 5.5.: Calltarget Classification

target	opt	problematic	-0	-1	-2	-3	-4	-5	-6
x	x	x	x	x	x	x	x	x	x

Table 5.6.: Coumpound table

target	opt	#	Callsites: param perf. %, probl %	#	Callsites: param perf. %, probl %
x	x	x	x	x	x

Table 5.7.: MAtching table

target	opt	ct	Ct probl.	at	at probl.	cs	clang cs probl.	padyn cs probl.
x	x	x	x	x	x	x	x	x

Table 5.8.: policy evaluation

target	opt	policy	0	1	2	3	4	5	6	sumarry
x	x	x	x	x	x	x	x	x	x	x

Table 5.9.: param wideness

5	4	3	2	1	0	param/wideness
x	x	x	x	x	x	0
x	x	x	x	x	x	8
x	x	x	x	x	x	16
x	x	x	x	x	x	32
x	x	x	x	x	x	64

Table 5.10.: tabelle 7

5	4	3	2	1	0	param/wideness
x	x	x	x	x	x	0
x	x	x	x	x	x	8
x	x	x	x	x	x	16
x	x	x	x	x	x	32
x	x	x	x	x	x	64

Table 5.11.: tabelle 7

5	4	3	2	1	0	param/wideness
x	x	x	x	x	x	0
x	x	x	x	x	x	8
x	x	x	x	x	x	16
x	x	x	x	x	x	32
x	x	x	x	x	x	64

alternative for abobe:

Figure 5.1.: impact of CFI and CFC




Table 5.12.: matching

target	opt	fn_count	fn_problem	at_count	at_problem	cs_count	cs_clang	cs_padyn
x	x	x	x	x	x	0	0	0

Figure 5.2.: liveness iteration, dummy




Figure 5.3.: reaching iteration, dummy

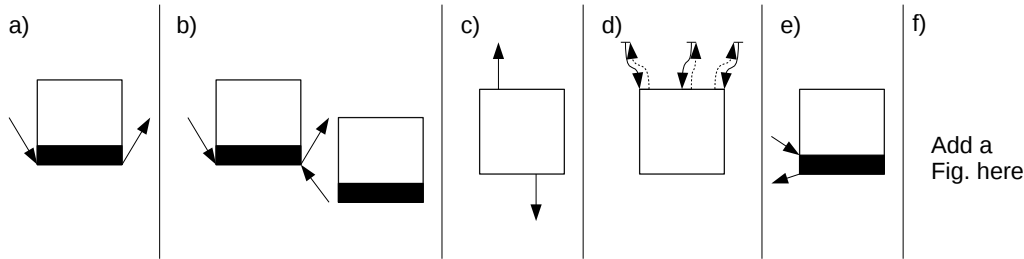


Table 5.13.: pairings compares

target	opt	policy	0	1	2	3	4	5	6	summary
proftpd	x	x	x	x	x	0	0	0	0	0
proftpd	x	x	x	x	x	0	0	0	0	0
vsftpd	x	x	x	x	x	0	0	0	0	0
vsftpd	x	x	x	x	x	0	0	0	0	0

Table 5.14.: policy baseline

target	opt	policy	0	1	2	3	4	5	6	summary
proftpd	x	x	x	x	x	0	0	0	0	0
proftpd	x	x	x	x	x	0	0	0	0	0
vsftpd	x	x	x	x	x	0	0	0	0	0
vsftpd	x	x	x	x	x	0	0	0	0	0

5.2. R2: Precision of our Tool

5.3. R3: Performance overhead of our Tool

5.4. R4: Instrumentation overhead of our Tool

5.5. R5: Security coverage of our tool

5.6. R6: Which kind of attacks can our tool defend off

5.7. R7: Whar are the limitations of our Tool

5.8. R8: To Do.

it is easier for the reader if we can directly map those section from underneath on the section from above.

5.9. Classification

5.9.1. Callsites

overestimation param count. table. number of parameters.

5.9.2. Calltargets

underestimation param table.

5.10. Patching Policies

Two types of diagrams. Table 5 from TypeArmor and a CDF to compare param count and param type. (baseline).

5.10.1. AT

5.10.2. ParamCount

table, cdf, baseline vs. server. approximations.

5.10.3. ParamType

table, cdf, baseline vs. server. approximations.

5.11. Security Evaluation

5.12. Performance

spec 2006.

6. Related Work

In this section we shortly review the main techniques and tools which are related to *Type-Shield*. Section 6.1 presents several techniques which can not fully defend against advanced CRAs. Section 6.2 gives an overview of tools used to recover type inference from binaries. Section 6.3 highlights several tools which can mitigate against the COOP attack on the binary level. Section 6.4 presents tools which can mitigate against the COOP source code level. Finally, section 6.5 shows some of the most promising runtime-based mitigation techniques against COOP.

6.1. Mitigation of Advanced Code-Reuse Attacks

COOP [3], Subversive-C [6] and Recursive-COOP [5] are advanced CRAs since these attacks: *i)* can not be addressed with shadow stacks techniques (i.e., do not violate the caller/callee convention), *ii)* coarse-grained CFI techniques are useless against these attacks, *iii)* hardware based approaches such as Intel CET [7] can not mitigate this attack for the same reason as in *i)*, and *iv)* OS-based approaches such as Windows Control Flow Guard [8] does not defend against COOP since the precomputed CFG does not contain edges for indirect call sites which are explicitly exploited during the COOP attack.

CRAs have many manifestations and it is out of scope of this work to list them all here. CRAs can be addressed in general the following ways: *(i)* binary instrumentation, *(ii)* source code recompilation and *(iii)* runtime application monitoring. While there is a plethora of tools and techniques which try to enforce CFI primitives in executables, source code and during runtime, we briefly list few of them in order for the reader to get familiar with the solution landscape. The approaches used to combat against CRAs are roughly based on the following techniques: *(a)* fine-grained CFI with hardware support, PathArmor [17], *(b)* coarse-grained CFI used for binary instrumentation, CCFIR [38], *(c)* coarse-grained CFI based on binary loader, CFCI [39] *(d)* fine-grained code randomization, O-CFI [36], *(e)* cryptography with hardware support, CCFI [35], *(f)* ROP stack pivoting, PBlocker [41], *(g)* canary based protection, DynaGuard [40], *(h)* checking vTable integrity for protecting against COOP based on CFI for source code such as SafeDispatch [2], vtv [4] LLVM and GCC compiler based vor vTable protection and binary rewriting such as vfGuard [42], vTint [37] and [5], *(i)* runtime and hardware support based on a combination of LBR, PMU and BTS registers CFIGuard [27] and *(j)* source code recompilation with CFI and/or randomization enforcement against JIT-ROP attacks, MCFI [26], RockJIT [30] and PiCFI [31].

Notice that the above techniques are useless against the aforementioned advanced CRAs, the list is not exhaustive and new protection techniques arise by combining available techniques or by using newly available hardware features.

6.2. Type-Inference on Executables

Recovering variable types from executable programs is very hard in general for several reasons. First, the quality of the disassembly can vary much from used framework to another. *TypeShield* is based on DynInst and the quality of the executable disassembly fits our needs. For a more comprehensive review on the capabilities of DynInst and other tools we advise the reader to have a look at [18].

Second, alias analysis in binaries is undecidable in theory and intractable in practice [23]. There are several most promising tools such as: Rewards [19], BAP [20], SmartDec [21], and Divine [22]. These tools try with more or less success to recover type information from binary programs with different goals. Typical goals are: *i*) full program reconstruction (binary to code conversion, reversing), *ii*) checking for buffer overflows, *iii*) integer overflows and other types of memory corruptions. For a more exhaustive review of such tools we advise the reader to have a look at the review of Caballero et al. [16]. Interesting to notice is that the code from only a few of these tools is available.

6.3. Binary-based Protection against COOP

TypeShield is most similar to TypeArmor [9] since we also enforce strong binary-level invariants on the number of function parameters. *TypeShield* similarly to TypeArmor targets exclusive protection against advanced exploitation techniques which can bypass fine-grained CFI schemes and VTable protections at the binary level.

However, *TypeShield* is much more precise than TypeArmor since its enforcing policy is also based on the types of the function parameters. This results in a more precise selection of caller/callee pairs on which the fine-grained CFI policy is enforced. Thus, we achieve a reduced runtime overhead than TypeArmor since we enforce our CFI policy on fewer caller/callee pairs than TypeArmor.

6.4. Source Code based Protection against COOP

There are several source code based tools which can successfully protect against the COOP attack. Such tools are: ShrinkWrap [10], IFCC/VTV [4], SafeDispatch [2], vTrust [13], Readactor++ [15], CPI [11] and the tool presented by Bounov et al. [12]. These tools profit from high precision since they have access to the full semantic context of the program though the scope of the compiler on which they are based. Because of this reason these tools target mostly other types of security problems than binary-based tools address. For example some last advances in compile based protection against code reuse attacks address mainly performance issues. Currently, most of the above presented tools are only forward edge enforcers of fine-grained CFI policies with an overhead from 1% up to 15%.

We are aware that there is still a long research path to go until binary based techniques can recuperate program based semantic information from executable with the same precision as compiler based tools. These paths could be even endless since compilers are optimized for speed and are designed to remove as much as possible semantic information from an executable in order to make the program run as fast as possible. In light of this fact, *TypeShield* is another attempt to recuperate just the needed semantic information (types

and number of function parameters from indirect call sites) in order to be able to enforce a precise and with low overhead primitive against COOP attacks.

Rather than claiming that the invariants offered by *TypeShield* are sufficient to mitigate all versions of the COOP attack we take a more conservative path by claiming that *TypeShield* further raises the bar w.r.t. what is possible when defending against COOP attacks on the binary level.

6.5. Runtime-based Protection against COOP

There is something available out there but I can not use it: anonymous. Long story short conclusion: There are several promising runtime-based line of defenses against advanced CRAs but none of them can successfully protect against the COOP attack.

IntelCET [7] is based on, ENDBRANCH, a new CPU instruction which can be used to enforce an efficient shadow stack mechanism. The shadow stack can be used to check during program execution if caller/return pairs match. Since the COOP attack reuses whole functions as gadgets and does not violate the caller/return convention than the new feature provided by Intel is useless in the face of this attack. Nevertheless other highly notorious CRAs may not be possible after this feature will be implemented mainstream in OSs and compilers.

Windows Control Flow Guard [8] is based on a user-space and kernel-space components which by working closely together can enforce an efficient fine-grained CFI policy based on a precomputed CFG. These new features available in Windows 10 can considerably raise the bar for future attacks but in our opinion advanced CRAs such as COOP are still possible due to the typical characteristics of COOP.

PathArmor [17] is yet another tool which is based on a precomputed CFG and on the LBR register which can give a string of 16 up to 32 pairs of from/to addresses of different types of indirect instructions such as `call`, `ret`, and `jump`. Because of the sporadic query of the LBR register (only during invocation of certain function calls) and because of the sheer amount of data which passes through the LBR register this approach has in our opinion a fair potential to catch different types of CRAs but we think that against COOP this tool can not be used. First, because of the fact that the precomputed CFG does not contain edges for all possible indirect call sites which are accessed during runtime and second, the LBR buffer can be easily tricked by adding legitimate indirect call sites during the COOP attack.

7. Discussion

We have to define which points make sense and then talk about each other
Suggestion:

7.1. How to make the type inference more precise?

1/2 page

7.2. Comparison with TypeArmor and why are we better than TypeArmor?

1/2 page

7.3. Whys is not TypeArmor working as it should to?

1/2 page

7.4. What is not clear bout TypeArmor?

1/2 page

7.5. What can for sure not work as in TypeArmor paper explained?

Furthermore, bla ...

8. Conclusion and Future Work

This section should be no longer than 2 pages. ideally exactly 2 pages would be sufficient.

8.1. Conclusion

In this research, we presented our tool X,

Specify the points you want to talk about. Write 2-3 sentences about each point.

Point 1

Point 2

Point 3

8.2. Future Work

In future we bla.

Point 1

Point 2

Point 3

Specify the points you want to talk about. Write 2-3 sentences about each point.

Bibliography

- [1] M. Zhang and R. Sekar, Control Flow Integrity for COTS Binaries, In *Proceedings of the 22nd USENIX conference on Security*, (USENIX Sec), ACM, pp. 337-352, 2013.
- [2] D. Jang et al., safeDispatch: Securing C++ Virtual Calls from Memory Corruption Attacks, In *Symposium on Network and Distributed System Security*, (NDSS), 2014.
- [3] Felix Schuster, Thomas Tendyck, Christopher Liebchen, Lucas Davi, Ahmad-Reza Sadeghi, Thorsten Holz, Counterfeit Object-oriented Programming, In *Proceedings of the 2015 IEEE Symposium on Security and Privacy*, (S&P), 2015.
- [4] C. Tice et al., Enforcing Forward-Edge Control-Flow Integrity in GCC and LLVM, In *Proceedings of the 23rd USENIX conference on Security*, (USENIX Sec), ACM, 2014.
- [5] S. Crane et al., It's a TRaP: Table Randomization and Protection against Function-Reuse Attacks, In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, (CCS), 2015.
- [6] Julian Lettner, Benjamin Kollenda, Andrei Homescu, Per Larsen, Felix Schuster, Lucas Davi, Ahmad-Reza Sadeghi, Thorsten Holz, Michael Franz, Subversive-C: Abusing and Protecting Dynamic Message Dispatch, In *USENIX Annual Technical Conference*, (ATC), 2016.
- [7] Intel, Intel CET, <http://blogs.intel.com/evangelists/2016/06/09/intel-release-new-technology-specifications-protect-rop-attacks/>, 2016.
- [8] Microsoft, Windows Control Flow Guard, [https://msdn.microsoft.com/en-us/library/windows/desktop/mt637065\(v=vs.85\).aspx](https://msdn.microsoft.com/en-us/library/windows/desktop/mt637065(v=vs.85).aspx), 2015.
- [9] Victor van der Veen, Enes Goktas, Moritz Contag, Andre Pawlowski, Xi Chen, Sanjay Rawat, Herbert Bos, Thorsten Holz, Elias Athanasopoulos, and Cristiano Giuffrida, A Tough call: Mitigating Advanced Code-Reuse Attacks At The Binary Level, In *Proceedings of the 2016 IEEE Symposium on Security and Privacy*, (S&P), 2016.
- [10] I. Haller, E. Goktas, E. Athanasopoulos, G. Portokalidis, H. Bos, ShrinkWrap: VTable Protection Without Loose Ends, In *Annual Computer Security Applications Conference*, (ACSAC), 2015.
- [11] Volodymyr Kuznetsov, László Szekeres, Mathias Payert, George Candea, R. Sekar, Dawn Song, Code-Pointer Integrity, In *Proceedings of the 11th USENIX Symposium on Operating Systems Design and Implementation*, (OSDI), 2014.

- [12] Dimitar Bounov, Rami Gökhan Kici, and Sorin Lerner, Protecting C++ Dynamic Dispatch Through VTable Interleaving, In *Symposium on Network and Distributed System Security*, (NDSS), 2016.
- [13] Chao Zhang, Scott A. Carr, Tongxin Li, Yu Ding, Chengyu Song, Mathias Payer and Dawn Song, VTrust: Regaining Trust on Virtual Calls, In *Symposium on Network and Distributed System Security*, (NDSS), 2016.
- [14] Stephen Crane, Stijn Volckaert, Felix Schuster, Christopher Liebchen, Per Larsen, Lucas Davi, Ahmad-Reza Sadeghi, Thorsten Holz, Bjorn De Sutter, Michael Franz, It's a TRaP: Table Randomization and Protection against Function-Reuse Attacks, In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, (CCS), 2015.
- [15] Stephen Crane, Stijn Volckaert, Felix Schuster, Christopher Liebchen, Per Larsen, Lucas Davi, Ahmad-Reza Sadeghi, Thorsten Holz, Bjorn De Sutter, Michael Franz, It's a TRaP: Table Randomization and Protection against Function-Reuse Attacks, In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, (CCS), 2015.
- [16] Juan Caballero, and Zhiqiang Lin, Type Inference on Executables, In *ACM Computing Surveys*, (CSUR), 2016.
- [17] V. van der Veen et al., Practical Context-Sensitive CFI, In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, (CCS), 2015.
- [18] D. Andriesse, X. Chen, V. van der Veen, A. Slowinska, and H. Bos, An In-Depth Analysis of Disassembly on Full-Scale x86/x64 Binaries, In *Proceedings of the USENIX Conference on Security*, (USENIX Sec), 2016.
- [19] Zhiqiang Lin Xiangyu Zhang Dongyan Xu, Automatic Reverse Engineering of Data Structures from Binary Execution, In *Symposium on Network and Distributed System Security*, (NDSS), 2010.
- [20] David Brumley, Ivan Jager, Thanassis Avgerinos, and Edward J. Schwartz, BAP: A Binary Analysis Platform, In *Proceedings of Computer Aided Verification*, (CAV), 2011.
- [21] Alexander Fokin, Yegor Derevenets, Alexander Chernov, and Katerina Troshina, SmartDec: Approaching C++ decompilation, In *Working Conference on Reverse Engineering*, (WCRE), 2011.
- [22] G. Balakrishnan and T. Reps, DIVINE: Discovering variables in executables, In *International Conference on Verification, Model Checking, and Abstract Interpretation*, (VMCAI), 2007.
- [23] Alan Mycroft, Lecture Notes, In <https://www.cl.cam.ac.uk/am21/papers/sas07slides.pdf>, 2007.
- [24] N. Carlini et al., Control-Flow Bending: On the Effectiveness of Control-Flow Integrity, In *Proceedings of the 24th USENIX conference on Security*, ACM, 2015.

- [25] R. Skowyra et al., Systematic Analysis of Defenses Against Return-Oriented Programming, In *Proceedings of the 16th International Symposium on Research in Attacks, Intrusions, and Defenses*, (RAID), 2013.
- [26] B. Niu et al., Modular Control-Flow Integrity, In *ACM Conference on Programming Language Design and Implementation*, (PLDI), 2014.
- [27] R. Skowyra et al., Hardware-Assisted Fine-Grained Code-Reuse Attack Detection, In *Proceedings of the 18th International Symposium on Research in Attacks, Intrusions, and Defenses*, (RAID), 2015.
- [28] E. Göktaş et al., Out Of Control: Overcoming Control-Flow Integrity, In *Proceedings of the 2014 IEEE Symposium on Security and Privacy*, (S&P), 2014.
- [29] M. Abadi et al., Control Flow Integrity, In *the 12th ACM Conference on Computer and Communications Security*, (CCS), 2005.
- [30] B. Niu et al., RockJIT: Securing Just-In-Time Compilation Using Modular Control-Flow Integrity, In *the 21th ACM Conference on Computer and Communications Security*, (CCS), 2014.
- [31] B. Niu et al., Per-Input Control-Flow Integrity, In *the 22th ACM Conference on Computer and Communications Security*, (CCS), 2015.
- [32] M. Abadi et al., Control Flow Integrity Principles, Implementations, and Applications, In *ACM Transactions on Information and System Security*, (TISSEC), 2009.
- [33] N. Carlini et al., ROP is still dangerous: Breaking Modern Defenses, In *Proceedings of the 23rd USENIX conference on Security*, (USENIX Sec), ACM, 2014.
- [34] M. Wang et al., Binary Code Continent: Finer-Grained Control Flow Integrity for Stripped Binaries, In *Annual Computer Security Applications Conference*, (ACSAC), 2015.
- [35] A. J. Mashtizadeh et al., CCFI: Cryptographically Enforced Control Flow Integrity, In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, (CCS), 2015.
- [36] V. Mohan et al., Opaque Control-Flow Integrity, In *Symposium on Network and Distributed System Security*, (NDSS), 2015.
- [37] C. Zhang et al., VTint: Protecting Virtual Function Tables Integrity, In *Symposium on Network and Distributed System Security*, (NDSS), 2015.
- [38] C. Zhang et al., Practical Control Flow Integrity & Randomization for Binary Executables, In *Proceedings of the 2013 IEEE Symposium on Security and Privacy*, (S&P), 2013.
- [39] M. Zhang et al., Control Flow and Code Integrity for COTS binaries: An Effective Defense Against Real-ROP Attacks, In *Annual Computer Security Applications Conference*, (ACSAC), 2015.
- [40] T. Petsios et al., DynaGuard: Armoring Canary-based Protections against Brute-force Attacks, In *Annual Computer Security Applications Conference*, (ACSAC), 2015.

- [41] A. Prakash et al., Defeating ROP Through Denial of Stack Pivot, In *Annual Computer Security Applications Conference, (ACSAC)*, 2015.
- [42] A. Prakash et al., Strict Protection for Virtual Function Calls in COTS C++ Binaries, In *Symposium on Network and Distributed System Security, (NDSS)*, 2015.